

Natural Language Processing & Word Embeddings

9/10 points (90.00%)

Quiz, 10 questions

✓ Congratulations! You passed![Next Item](#)1 / 1
points

1.

Suppose you learn a word embedding for a vocabulary of 10000 words. Then the embedding vectors should be 10000 dimensional, so as to capture the full range of variation and meaning in those words.

☐ True☒ False**Correct**

The dimension of word vectors is usually smaller than the size of the vocabulary. Most common sizes for word vectors ranges between 50 and 400.

1 / 1
points

2.

What is t-SNE?

☐ A linear transformation that allows us to solve analogies on word vectors☒ A non-linear dimensionality reduction technique**Correct**

Yes

☐ A supervised learning algorithm for learning word embeddings☐ An open-source sequence modeling library

1 / 1
points

Natural Language Processing & Word Embeddings

9/10 points (90.00%)

Quiz, 10 questions

Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You then use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set.

x (input text)	y (happy?)
I'm feeling wonderful today!	1
I'm bummed my cat is ill.	0
Really enjoying this!	1

Then even if the word "ecstatic" does not appear in your small training set, your RNN might reasonably be expected to recognize "I'm ecstatic" as deserving a label $y = 1$.



True

Correct

Yes, word vectors empower your model with an incredible ability to generalize. The vector for "ecstatic" would contain a positive/happy connotation which will probably make your model classified the sentence as a "1".



False

1 / 1
points

4.

Which of these equations do you think should hold for a good word embedding? (Check all that apply)



$$e_{boy} - e_{girl} \approx e_{brother} - e_{sister}$$

Correct

Yes!



$$e_{boy} - e_{girl} \approx e_{sister} - e_{brother}$$

Un-selected is correct

Natural Language Processing & Word Embeddings

9/10 points (90.00%)

Quiz, 10 questions

Correct

Yes!



$$e_{boy} - e_{brother} \approx e_{sister} - e_{girl}$$

Un-selected is correct0 / 1
points

5.

Let E be an embedding matrix, and let o_{1234} be a one-hot vector corresponding to word 1234. Then to get the embedding of word 1234, why don't we call $E * o_{1234}$ in Python?



It is computationally wasteful.

The correct formula is $E^T * o_{1234}$.

This doesn't handle unknown words (<UNK>).



None of the above: calling the Python snippet as described above is fine.

This should not be selected

No, one of the above is true.

1 / 1
points

6.

When learning word embeddings, we create an artificial task of estimating $P(\text{target} \mid \text{context})$. It is okay if we do poorly on this artificial prediction task; the more important by-product of this task is that we learn a useful set of word embeddings.



True

Correct



Natural Language Processing & Word Embeddings

9/10 points (90.00%)

Quiz, 10 questions



1 / 1
points

7.

In the word2vec algorithm, you estimate $P(t | c)$, where t is the target word and c is a context word. How are t and c chosen from the training set? Pick the best answer.

- ☐ c is the sequence of all the words in the sentence before t .
- ☐ c is the one word that comes immediately before t .
- ☐ c is a sequence of several words immediately before t .
- ☒ c and t are chosen to be nearby words.

Correct



1 / 1
points

8.

Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The word2vec model uses the following softmax function:

$$P(t | c) = \frac{e^{\theta_t^T e_c}}{\sum_{t'=1}^{10000} e^{\theta_{t'}^T e_c}}$$

Which of these statements are correct? Check all that apply.



θ_t and e_c are both 500 dimensional vectors.

Correct



θ_t and e_c are both 10000 dimensional vectors.

Un-selected is correct



θ_t and e_c are both trained with an optimization algorithm such as Adam or gradient descent.

Correct

Natural Language Processing & Word Embeddings

9/10 points (90.00%)

Quiz, 10 questions



After training, we should expect θ_t to be very close to e_c when t and c are the same word.

**Un-selected is correct**1 / 1
points

9.

Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The GloVe model minimizes this objective:

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij})(\theta_i^T e_j + b_i + b'_j - \log X_{ij})^2$$

Which of these statements are correct? Check all that apply.



θ_i and e_j should be initialized to 0 at the beginning of training.

**Un-selected is correct**

θ_i and e_j should be initialized randomly at the beginning of training.

**Correct**

X_{ij} is the number of times word i appears in the context of word j .

**Correct**

The weighting function $f(\cdot)$ must satisfy $f(0) = 0$.

**Correct**

The weighting function helps prevent learning only from extremely common word pairs. It is not necessary that it satisfies this function.

1 / 1
points

10.

You have trained word embeddings using a text dataset of m_1 words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of m_2 words. Keeping in mind that using word embeddings is a form of transfer learning, under which of these circumstances would you expect the word embeddings to be helpful?

 $m_1 \gg m_2$ **Correct** $m_1 \ll m_2$

1

